

R: Please re-write the following expression using the pipe (%>%) operator.

```
mean(is.na(age))
```

Session 4 - dplyr

R: suppose you have a data frame `patients` with the following contents:

	name	age	height	weight
1	Bill	25	181	105
2	Gina	28	160	60
3	Kelly	26	171	75
4	Sean	38	168	80

Assume you are only interested in the `height` and `weight` columns.

Using the pipe operator and appropriate dplyr functions, create a data frame `heights_weights` which contains the `height` and `weight` columns of `patients`.

R: suppose you have a data frame `patients` with the following contents:

	name	age	height	weight
1	Bill	25	181	105
2	Gina	28	160	60
3	Kelly	26	171	75
4	Sean	38	168	80

Assume you'd like to de-identify the data set.

Using the pipe operator and appropriate dplyr functions, create a data frame `patients_deidentified` which contains all the same columns as `patients` except `name`.

R: suppose you have a data frame `patients` with the following contents:

	name	age	height	weight
1	Bill	25	181	105
2	Gina	28	160	60
3	Kelly	26	171	75
4	Sean	38	168	80

Assume you're only interested in patients younger than 30 years old.

Using the pipe operator and appropriate dplyr functions, create a data frame `patients_less_than_30` which contains only the rows of `patients` whose age is less than 30.

R: suppose you have a data frame `patients` with the following contents:

	name	age	height	weight
1	Bill	25	181	105
2	Gina	28	160	60
3	Kelly	26	171	75
4	Sean	38	168	80

Assume you're interested in patients by the name Bill or Sean.

Using the pipe operator and appropriate dplyr functions, create a data frame `bill_or_sean` which contains only the rows of `patients` whose name is either Bill OR Sean.

R: suppose you have a data frame `patients` with the following contents:

	name	age	height	weight
1	Bill	25	181	105
2	Gina	28	160	60
3	Kelly	26	171	75
4	Sean	38	168	80

Assume you'd like to sort the patients by age.

Using the pipe operator and appropriate dplyr functions, create a data frame `patients_by_age` which contains the same data as `patients` but is sorted in ascending order, by age.

R: suppose you have a data frame `patients` with the following contents:

	name	age	height	weight
1	Bill	25	181	105
2	Gina	28	160	60
3	Kelly	26	171	75
4	Sean	38	168	80

Assume you'd like to add column containing the body mass index (BMI) of each patient. The BMI is the weight (in kg) divided by the height (in m) squared.

Using the pipe operator and appropriate dplyr functions, modify the data frame `patients` to add a column `bmi` which is equal to `weight / (height / 100) ^ 2`

R: suppose you have a data frame `patients` with the following contents:

	name	age	sex	height	weight	bmi
1	Bill	25	male	181	105	32.1
2	Gina	28	female	160	60	23.4
3	Kelly	26	female	171	75	25.6
4	Sean	38	male	168	80	28.3

Let's say we're interested in the mean and standard deviation of the BMI of the patients, broken down by sex.

Using the pipe operator and appropriate dplyr functions, do the following:

1. Group the `patients` data by `sex`, then
2. Create a data frame with the following summary statistics:
 1. `n`: the number of cases in each group
 2. `bmi_mean`: the mean of `bmi` of each group
 3. `bmi_sd`: the standard deviation of `bmi` of each group

The resulting table should look like this:

	sex	n	bmi_mean	bmi_sd
1	female	2	24.5	1.555635
2	male	2	30.2	2.687006

R: suppose you have the following two data frames:

	> bmi_data				> patients	
	id	age	sex	bmi	id	name
2	28	female	23.4		1	Bill
3	26	female	25.6		2	Gina
4	38	male	28.3		3	Kelly
1	25	male	32.1		4	Sean

Let's say we want to augment `bmi_data` with a column `name` containing the patient name, identified by the column `id` that is present in both tables.

Join `bmi_data` and `patients` to add a new column `name` to `bmi_data`, using `id` as the key to join on. Use the pipe operator where appropriate. Overwrite `bmi_data` with the result.